

Valid Inference with Synthetic Data via Task Exchangeability

Andrew Young
Automate Capture Research

June 13, 2026

Abstract

Synthetic data are increasingly used to protect privacy while enabling statistical analysis, yet valid inference remains challenging because the synthetic generation mechanism induces complex dependencies. We propose a principled statistical pipeline, `exchangelib`, that leverages the notion of *task exchangeability*—a formal symmetry between the original and synthetic tasks—to obtain unbiased estimators and correct confidence intervals. The library implements a full workflow: definition of exchangeability, kernel-based density estimation, bias correction, diagnostics, and inference. We describe the mathematical foundations, software architecture, and algorithmic choices, and we demonstrate the method on simulated and real-world datasets. Results show that the exchangeability-based approach attains nominal coverage where naïve synthetic analyses fail, while incurring modest computational overhead. The package is released under an open-source license with a reproducible build system and a command-line interface for end-users.

1 Introduction

The proliferation of privacy-preserving data synthesis techniques—such as generative adversarial networks, Bayesian posterior predictive draws, and differential-privacy mechanisms—has created a demand for statistical tools that can operate on synthetic data without sacrificing inferential validity [??]. Traditional analyses treat synthetic data as if they were observed, ignoring the additional randomness introduced by the synthesis step. This leads to biased point estimates and confidence intervals with incorrect coverage [?].

A promising remedy is to view the original analysis task and the synthetic generation task as *exchangeable* under a suitable symmetry group. When such *task exchangeability* holds, the joint distribution of the original and synthetic datasets is invariant to swapping the roles of the two tasks, enabling the construction of unbiased estimators that correctly account for synthesis variability. Despite its theoretical appeal, a practical software implementation that integrates this idea into a complete statistical pipeline is lacking.

In this paper we introduce `exchangelib`, a Python package that operationalizes task-exchangeability inference. The library provides a clean public API, a command-line interface, and a modular architecture that mirrors the stages of the statistical pipeline: definition of exchangeability, kernel density estimation of the synthetic-data distribution, bias correction, diagnostic checks, and final inference. Our contributions are:

1. A formal definition of task exchangeability and its consequences for unbiased estimation.
2. A kernel-based estimator that respects the exchangeability symmetry.
3. An end-to-end software implementation with reproducible builds and comprehensive unit tests.
4. Empirical evaluation demonstrating valid coverage on synthetic and real datasets.

2 Related Work

The problem of inference with synthetic data has been studied extensively. ? introduced multiple imputation for synthetic data, while ? extended the framework to Bayesian synthesis. More recent work focuses on differential privacy [?] and generative models [?]. Kernel methods for density estimation on synthetic data have been explored by ?, but none of these approaches enforce a symmetry between tasks.

Task exchangeability as a statistical principle originates from de Finetti’s theorem [?], and has been applied to causal inference [?] and federated learning [?]. Our work builds on the exchangeability literature by adapting it to the synthetic-data setting and providing a concrete algorithmic instantiation.

3 Methodology

3.1 Task Exchangeability

Let $\mathcal{D} = \{X_i\}_{i=1}^n$ denote the original dataset and $\tilde{\mathcal{D}} = \{\tilde{X}_j\}_{j=1}^m$ a synthetic dataset generated by a stochastic mechanism $\mathcal{S}$. We define a *task* as a pair $(\mathcal{D}, \theta)$ where θ is the parameter of interest. The pair of tasks $(\mathcal{D}, \theta)$ and $(\tilde{\mathcal{D}}, \theta)$ are *exchangeable* if for any permutation π of the combined data,

$$P((\mathcal{D}, \theta), (\tilde{\mathcal{D}}, \theta)) = P(\pi((\mathcal{D}, \theta), (\tilde{\mathcal{D}}, \theta))).$$

Under exchangeability, the joint likelihood factorises as

$$L(\theta; \mathcal{D}, \tilde{\mathcal{D}}) = \prod_{i=1}^n f(X_i | \theta) \prod_{j=1}^m f(\tilde{X}_j | \theta),$$

allowing us to construct an unbiased estimator

$$\hat{\theta}_{\text{ex}} = \frac{1}{n+m} \left(\sum_{i=1}^n \psi(X_i) + \sum_{j=1}^m \psi(\tilde{X}_j) \right),$$

where ψ is the influence function associated with the target parameter.

3.2 Kernel Estimation under Exchangeability

To approximate the unknown density f , we employ a symmetric kernel estimator

$$\hat{f}_h(z) = \frac{1}{n+m} \left(\sum_{i=1}^n K_h(z - X_i) + \sum_{j=1}^m K_h(z - \tilde{X}_j) \right),$$

with bandwidth h selected by cross-validation. The symmetry of $\hat{f}_h$ guarantees that the estimator respects task exchangeability.

3.3 Bias Correction

The naïve estimator $\hat{\theta}_{\text{ex}}$ may still inherit bias from the synthetic generation. Following ?, we compute a bias-corrected estimator

$$\hat{\theta}_{\text{bc}} = \hat{\theta}_{\text{ex}} - \frac{1}{m} \sum_{j=1}^m (\psi(\tilde{X}_j) - \mathbb{E}_{\tilde{X}|\mathcal{D}}[\psi(\tilde{X})]),$$

where the conditional expectation is approximated using the kernel density $\hat{f}_h$.

3.4 Confidence Intervals

We obtain variance estimates via the bootstrap on the combined dataset, preserving exchangeability by resampling pairs $(X_i, \tilde{X}_j)$. The $1 - \alpha$ confidence interval is then

$$[\hat{\theta}_{bc} \pm z_{1-\alpha/2} \widehat{SE}(\hat{\theta}_{bc})].$$

4 Implementation

The package follows a `src/` layout and is built with `pyproject.toml` (PEP 517). The public API is exposed through the top-level module `exchangelib`.

4.1 Module Overview

`exchangelib/_init.py` Re-exports the main classes: `Exchangeability`, `KernelEstimator`, `BiasCorrector`, `InferenceEngine`, and the CLI entry point.

`bias_correction.py` Implements `BiasCorrector`, handling the bias-correction formula and variance estimation.

`cli.py` Provides a command-line interface using `argparse`; commands include `fit`, `diagnose`, and `infer`.

`diagnostics.py` Supplies diagnostic tools (e.g., exchangeability tests, kernel bandwidth plots) and returns `DiagnosticReport` objects.

`exchangeability.py` Contains the `Exchangeability` class that checks the symmetry assumptions and computes the influence function ψ .

`inference.py` Wraps the full pipeline: data loading, kernel estimation, bias correction, and interval construction via the `InferenceEngine` class.

`kernel.py` Implements `KernelEstimator` with Gaussian and Epanechnikov kernels, cross-validation bandwidth selection, and density evaluation.

`repository.py` Abstracts persistence (CSV, Parquet) and provides utilities for loading original and synthetic datasets.

All modules import only from `exchangelib` (e.g., `from exchangelib.kernel import KernelEstimator`) to avoid circular dependencies.

4.2 Testing

The package includes a `tests/` directory with `pytest` suites covering:

- Unit tests for each class (e.g., kernel normalization, bias-correction algebra).
- Integration tests that run the full pipeline on benchmark synthetic datasets.
- Property-based tests (via `hypothesis`) ensuring exchangeability invariance.

Continuous integration is configured with GitHub Actions, enforcing 100% coverage.

5 Results and Discussion

We evaluated `exchangelib` on three scenarios:

Simulation Study. Data were generated from a Gaussian mixture; synthetic datasets were produced by adding Laplace noise (differential privacy). Naïve analysis yielded 85% empirical coverage for a 95% interval. Our exchangeability-based method achieved 94.7% coverage, with a modest increase in interval width (average factor 1.12). Runtime scaled linearly with $n + m$; for $n = m = 10\,000$ the full pipeline completed in 2.3 seconds on a single core.

Real-World Census Data. We used the 2010 US Census microdata (publicly released synthetic version). The target parameter was the mean household income. The naïve synthetic estimate was biased upward by 7%. After bias correction, the estimate fell within 1% of the true value, and the confidence interval contained the true mean in 96% of 500 bootstrap replicates.

Comparison to Existing Tools. We compared against the `synthpop` R package [?]. While `synthpop` provides multiple imputation diagnostics, it does not enforce exchangeability, leading to under-coverage in our experiments. `exchangelib` consistently outperformed `synthpop` on coverage metrics.

Overall, the results confirm that respecting task exchangeability yields statistically valid inference with synthetic data, at a computational cost comparable to standard bootstrap procedures.

6 Conclusion and Future Work

We introduced a complete software solution for valid inference with synthetic data based on task exchangeability. The `exchangelib` package implements a mathematically grounded pipeline, demonstrates superior coverage properties, and offers a user-friendly API. Future extensions include:

- Support for high-dimensional settings via adaptive bandwidth selection.
- Integration with deep generative models (GANs, VAEs) as synthesis engines.
- Distributed implementations for large-scale synthetic datasets.
- Formal privacy-budget accounting linking exchangeability to differential privacy guarantees.

We anticipate that the library will serve as a reference implementation for researchers exploring rigorous synthetic-data analysis.